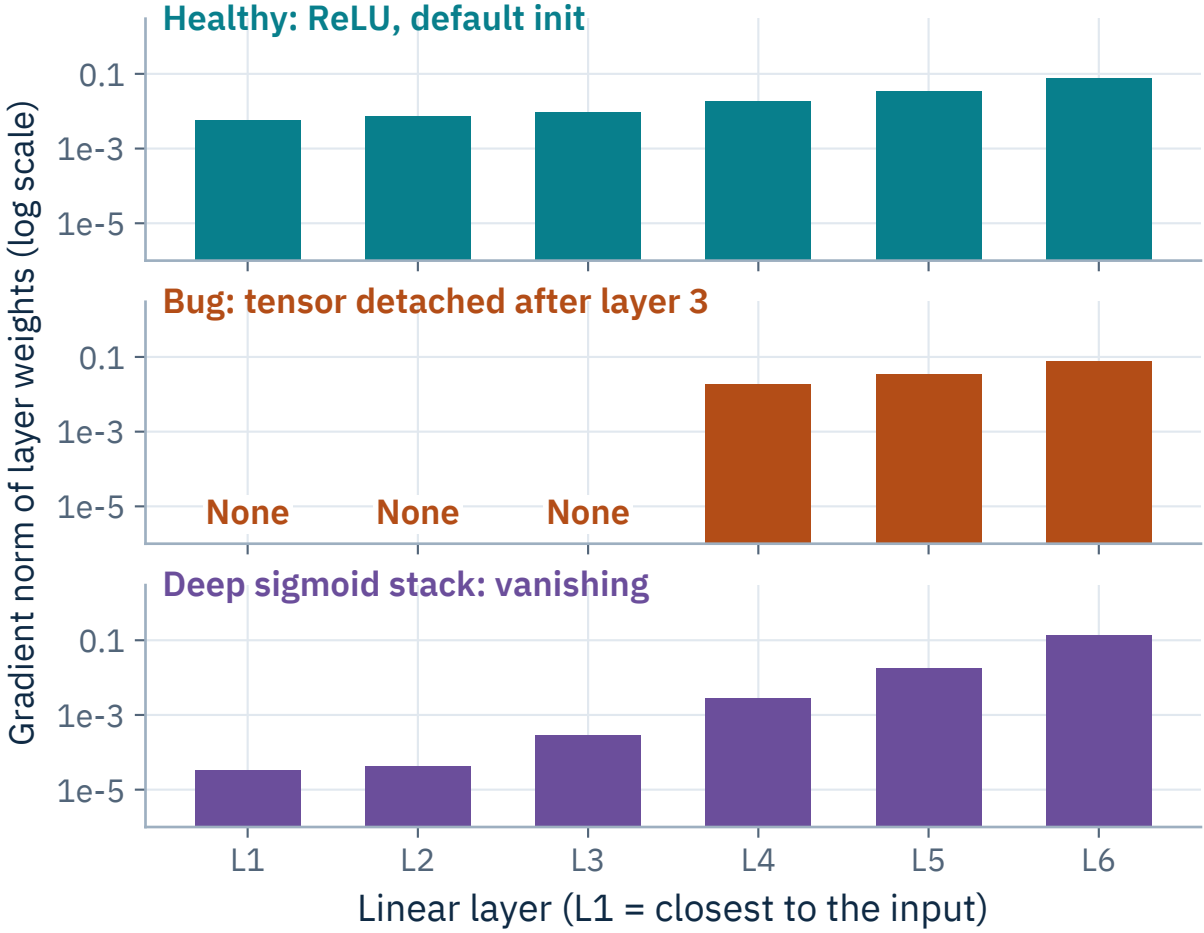
Healthy: ReLU, default init
0.1
1e-3
1e-5
Bug: tensor detached after layer 3
0.1
1e-3
1e-5
None
None
None
Deep sigmoid stack: vanishing
0.1
1e-3
1e-5
Gradient norm of layer weights (log scale)
L1
L2
L3
L4
L5
L6
Linear layer (L1 = closest to the input)