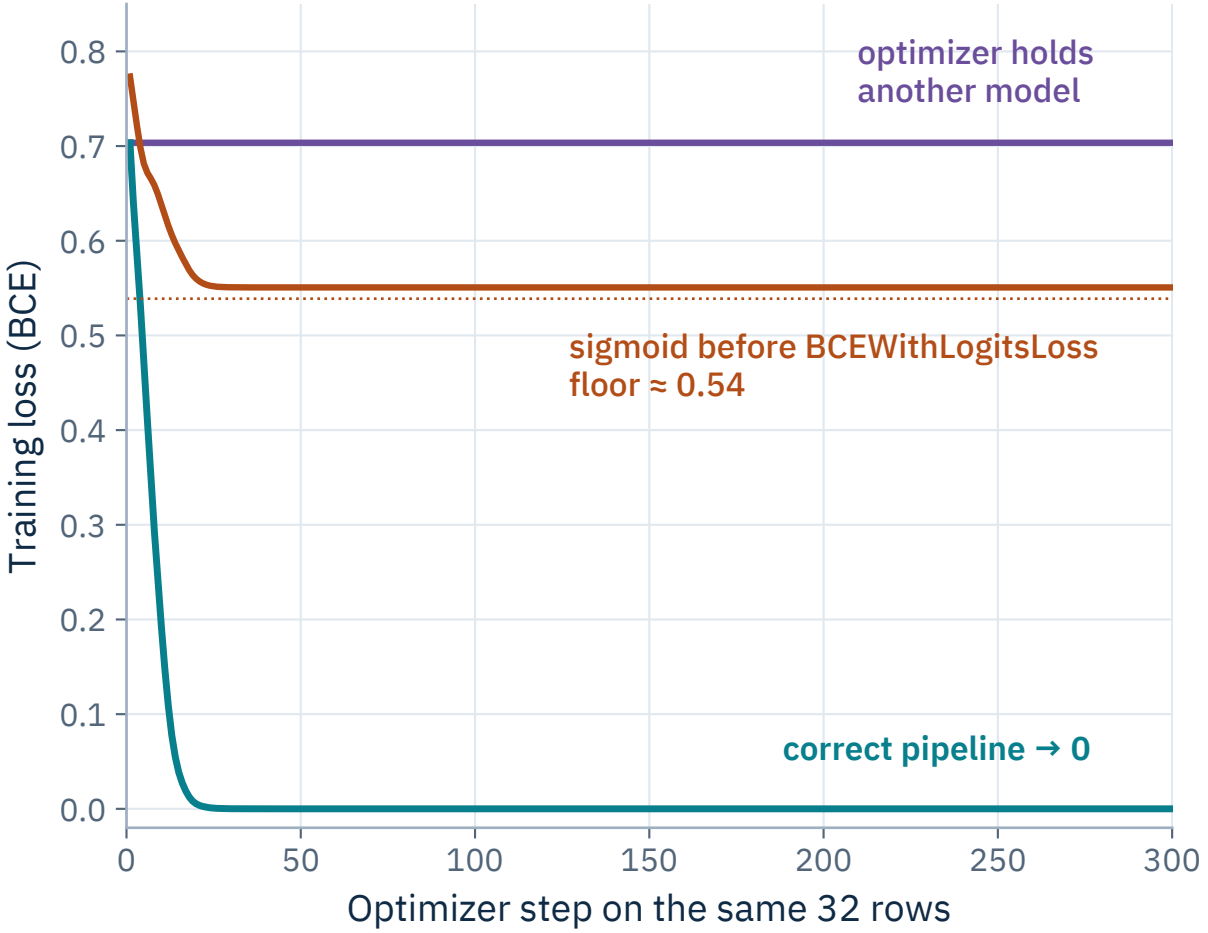
optimizer holds another model
sigmoid before BCEWithLogitsLoss
floor ≈ 0.54
correct pipeline → 0
Training loss (BCE)
0.8
0.7
0.6
0.5
0.4
0.3
0.2
0.1
0.0
0
50
100
150
200
250
300
Optimizer step on the same 32 rows