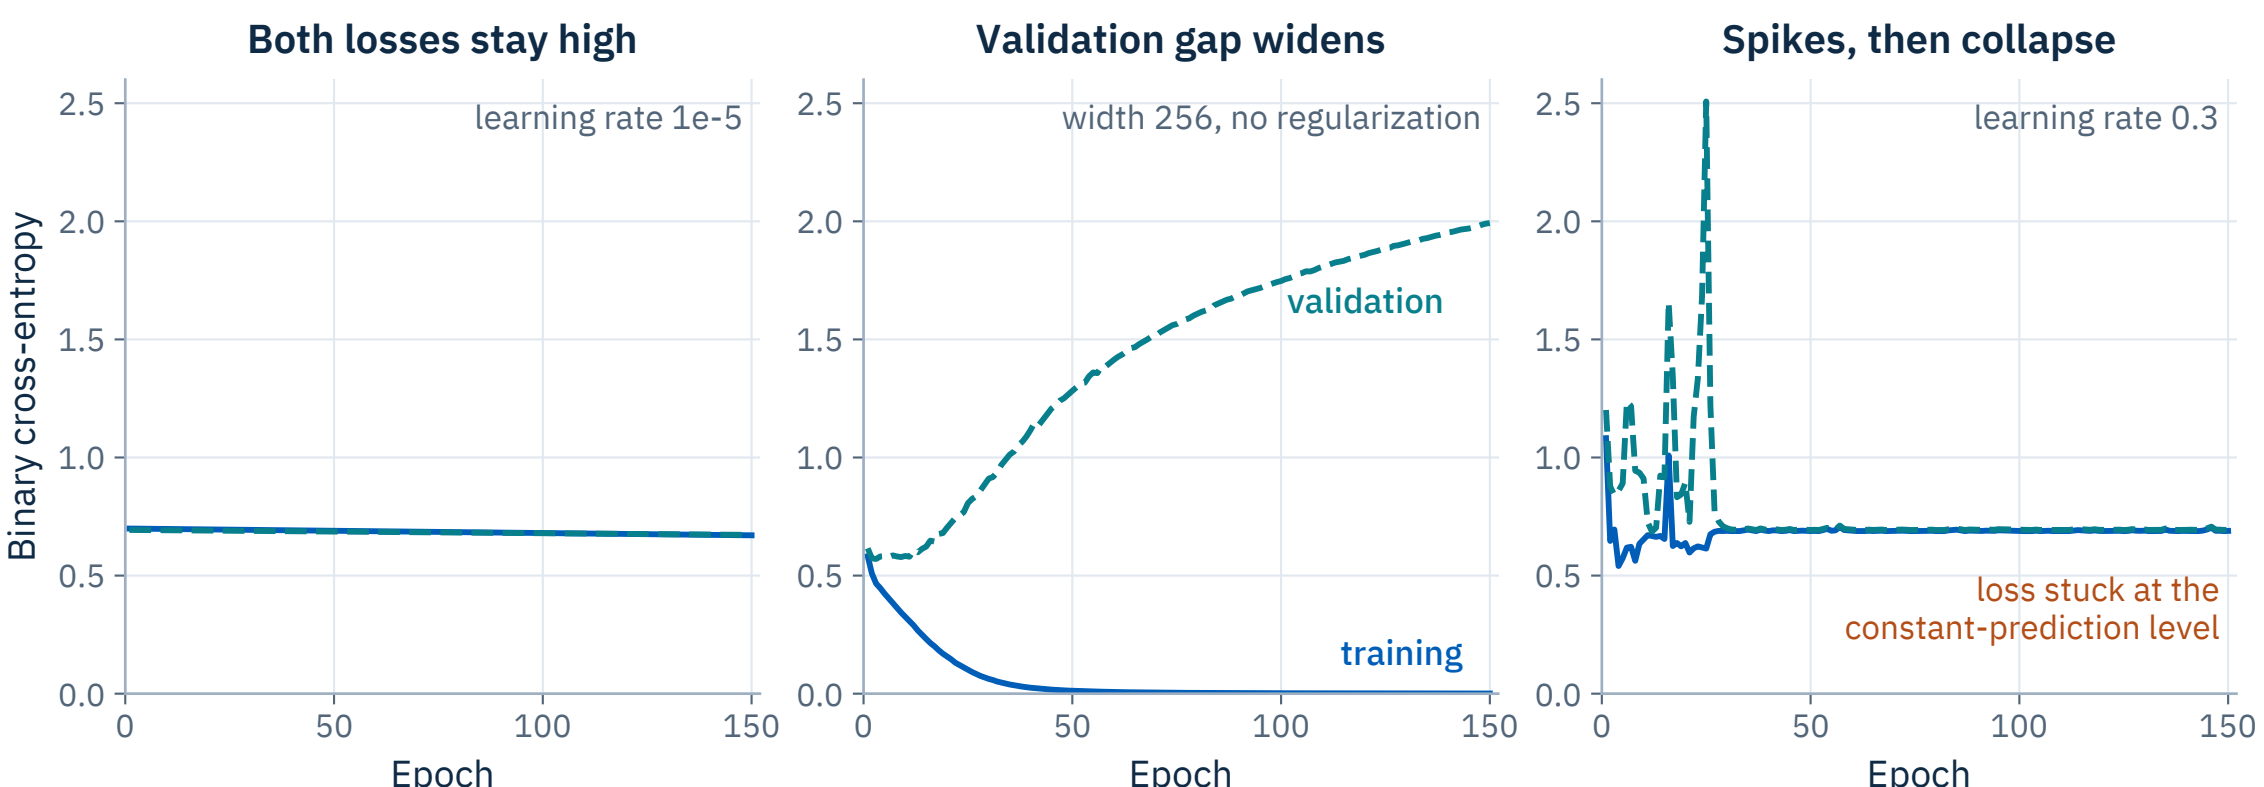
Both losses stay high
learning rate 1e-5
Validation gap widens
width 256, no regularization
validation
training
Spikes, then collapse
learning rate 0.3
loss stuck at the
constant-prediction level
Binary cross-entropy
Epoch
Epoch
Epoch
0.0
0.5
1.0
1.5
2.0
2.5
0
50
100
150