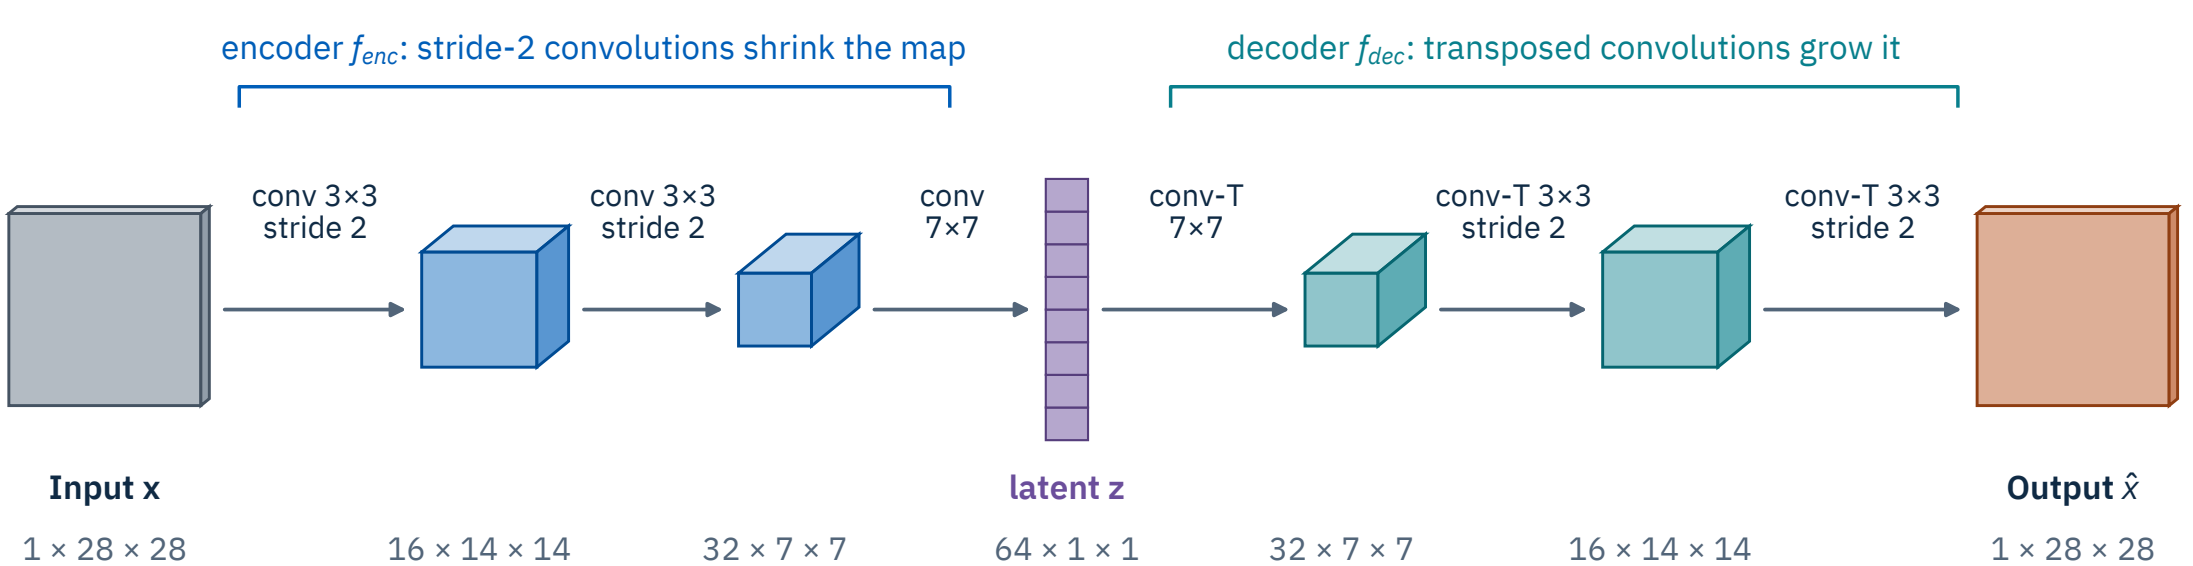
encoder f_{enc}: stride-2 convolutions shrink the map
decoder f_{dec}: transposed convolutions grow it
conv 3×3 stride 2
conv 3×3 stride 2
conv 7×7
conv-T 7×7
conv-T 3×3 stride 2
conv-T 3×3 stride 2
Input x
1 × 28 × 28
16 × 14 × 14
32 × 7 × 7
latent z
64 × 1 × 1
32 × 7 × 7
16 × 14 × 14
Output $\hat{x}$
1 × 28 × 28
training loss: MSE between $\hat{x}$ and the clean target x