Max pool 2 × 2, stride 2

fixed rule, 0 parameters

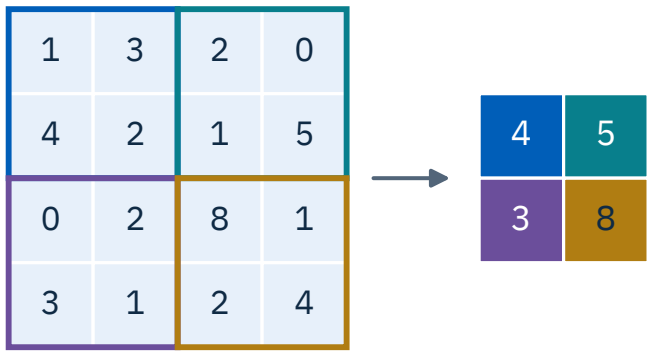

keeps the largest value per window

Strided convolution 2 × 2, stride 2

learned: 4 weights + 1 bias

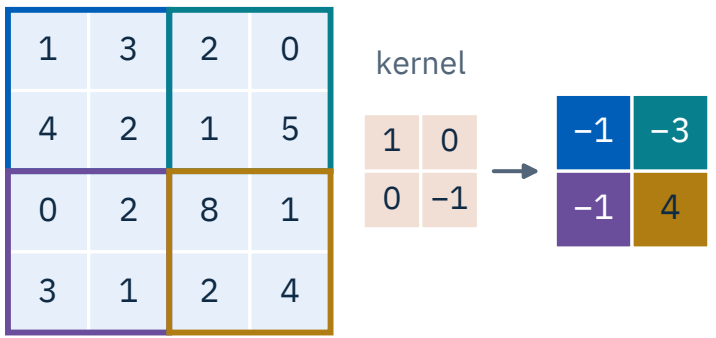

weighted sum per window, learned

Global average pooling

one number per channel

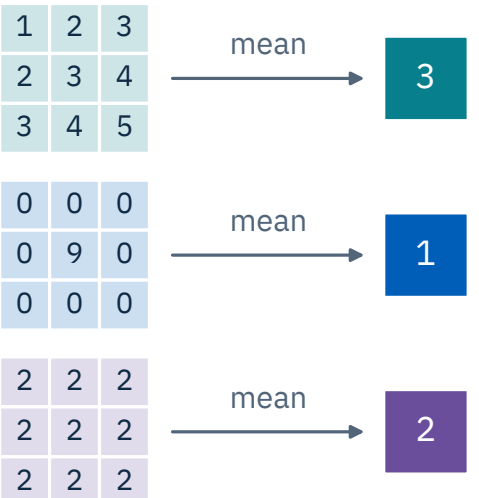

C × H × W → C, for any H × W